
The Model & Basic Computations

Chapter 1 and 2

The Model

Broadcast

Spanning Tree Construction

Traversal


Wake-up

Spanning Tree Construction

A spanning tree T of a graph $G = (V, E)$ is an acyclic subgraph of G such that $T = (V, E')$ and $E' \subseteq E$.


Assumptions:

single initiator
bidirectional links
total reliability
 G connected




Protocol SHOUT

Initially: $\forall x, \text{Tree-neighbors}(x) = \{ \}$



At the end:



$\forall x, \text{Tree-neighbors}(x) = \{\text{links that belong to the spanning tree}\}$



$\text{Tree-neighbors}(x)$ is a subset of $N(x)$

1.


init




Q? = do you want to be
my neighbour
in the spanning tree ?

2.


Q?



If it is the first time:



If I have already answered
yes to someone else:



States $S = \{\text{INITIATOR}, \text{IDLE}, \text{ACTIVE}, \text{DONE}\}$

$S_{\text{init}} = \{\text{INITIATOR}, \text{IDLE}\}$

$S_{\text{term}} = \{\text{DONE}\}$

Restrictions = $R; UI$

$R = \{\text{BL}, \text{TR}, \text{CN}\}$

INITIATOR

Spontaneously

root := true

Tree-neighbours := { }

send(Q) to N(x)

counter := 0

become ACTIVE

IDLE

receiving(Q)

root := false

parent := sender

Tree-neighbours := {sender}

send(yes) to sender

counter := 1

if counter = |N(x)| then

 become DONE

else

 send(Q) to N(x) - {sender}

 become ACTIVE

ACTIVE

receiving(Q)

send(no) to sender

receiving(yes)

Tree-neighbours:=

Tree-neighbours U sender

counter := counter +1

if counter = $|N(x)|$

become DONE

receiving(no)

counter := counter +1

if counter = $|N(x)|$

become DONE

Note:

SHOUT = FLOOD + REPLY

Correctness and Termination

- If x is in Tree-neighbours of y , y is in Tree-neighbours of x
- If x sends YES to y , then x is in Tree-neighbour of y and is connected to the initiator by a chain of YES
- Every x (except the initiator) sends exactly one YES



The spanning graph defined by the Tree-neighbour relation is connected, acyclic and contains all the entities

Note: local termination



Message Complexity

$SHOUT = FLOOD + REPLY$




$Messages(SHOUT) = 2 M(FLOOD)$

Possible situations




Impossible situations




Message Complexity - worst case

Total n. of Q:





only one Q on the ST links



on the other links

$$\begin{aligned}\text{Total: } & 2(m - (n-1)) + (n-1) \\ & = 2m - n + 1\end{aligned}$$

Message Complexity - worst case



Message Complexity - worst case

$$\begin{aligned} & 2m - n + 1 + 2(m - (n-1)) + n-1 \\ &= 2m - n + 1 + 2m - 2n + 2 + n - 1 \\ &= 4m - 2n + 2 \end{aligned}$$

$$\text{Messages}(\text{SHOUT}) = 4m - 2n + 2$$

In fact: $M(\text{SHOUT}) = 2 M(\text{FLOOD}) = 2(2m - n + 1)$

$\Omega(m)$ is a lower bound also in this case
therefore SHOUT is optimal asymptotically

Spanning Tree Construction

Without "NO"

Protocol SHOUT+

States $S = \{\text{INITIATOR}, \text{IDLE}, \text{ACTIVE}, \text{DONE}\}$

$S_{\text{init}} = \{\text{INITIATOR}, \text{IDLE}\}$

$S_{\text{term}} = \{\text{DONE}\}$

INITIATOR

Spontaneously

root := true

Tree-neighbours := { }

send(Q) to N(x)

counter := 0

become ACTIVE

IDLE

receiving(Q)

root := false

parent := sender

Tree-neighbours := {sender}

send(yes) to sender

counter := 1

if counter = |N(x)| then

become DONE

else

send(Q) to N(x) - {sender}

become ACTIVE

ACTIVE

receiving(Q) (to be interpreted as NO)

counter := counter + 1
if counter = $|N(x)|$
become DONE

receiving(yes)


Tree-neighbours :=
Tree-neighbours \cup {sender}
counter := counter + 1
if counter = $|N(x)|$
become DONE

On each link there will be exactly 2 messages:



either

or



$$\text{Messages}(\text{SHOUT}+) = 2m$$

Much better than:

$$\text{Messages}(\text{SHOUT}) = 4m - 2n + 2$$

Spanning Tree Construction

With Notification

States $S = \{\text{INITIATOR}, \text{IDLE}, \text{ACTIVE}, \text{DONE}\}$

$S_{\text{init}} = \{\text{INITIATOR}, \text{IDLE}\}$

$S_{\text{term}} = \{\text{DONE}\}$

INITIATOR

Spontaneously

root := true

Tree-neighbours := { }

send(Q) to N(x)

counter := 0

ack-counter := 0

become ACTIVE

IDLE

receiving(Q)

root := false

parent := sender

Tree-neighbours := {sender}

send(yes) to sender

counter := 1

ack-counter := 0

if counter = |N(x)| then

CHECK

else

send(Q) to N(x) - {sender}

become ACTIVE

ACTIVE

receiving(Q)

counter := counter +1

if counter = $|N(x)|$ and not root then

CHECK

receiving(yes)

Tree-neighbours := Tree-neighbours \cup {sender}

counter := counter +1

if counter = $|N(x)|$ and not root then

CHECK

ACTIVE (cont)

receiving(Ack)

ack-counter := ack-counter + 1

if counter = $|N(x)|$ // indicate tree-neighbors is done

if root then

if ack-counter = $|Tree-neighbours|$

send(Terminate) to Tree-neighbours

become DONE

else if ack-counter = $|Tree-neighbours| - 1$

send(Ack) to parent

receiving(Terminate)

send(Terminate) to Children

become DONE

// Children is $Tree-neighbours - \{parent\}$





CHECK

If I am a leaf

(* that is: Children := Tree-neighbours - {parent}
if Children = emptyset *)

send(Ack) to parent

What happens if there are multiple initiators ?



Impossibility Result

Theorem

The SPT problem is deterministically unsolvable under \mathcal{R}

no deterministic protocol that always correctly terminates within finite time

The idea:

- all initiator entities start in the same state and maintain it for the whole execution
- a correct solution requires entities end up with different states

An election is needed to have a unique initiator.

or

Another protocol has to be devised.

NOTE: Election is impossible if the nodes do not have distinct IDs

Traversal

Depth First Search

Assumptions

- Single initiator
- Bidirectional links
- No faults
- G connected

$S = \{\text{INITIATOR}, \text{VISITED}, \text{DONE}\}$

One version

- 1) When first visited, remember who sent,
forward the token to one of the unvisited neighbours
wait for its reply
- 2) When neighbour receives,
if already visited, it will return the token saying it is a
back edge
otherwise, will forward it (sequentially)
to all its unvisited neighbour before returning it
- 3) If there are no more unvisited neighbours, return the token (reply)
to the node from which it first received the token
- 4) Upon reception of reply, forward the token to another
unvisited neighbour

States $S = \{\text{INITIATOR}, \text{IDLE}, \text{VISITED}, \text{DONE}\}$

$S_{\text{init}} = \{\text{INITIATOR}, \text{IDLE}\}$

$S_{\text{term}} = \{\text{DONE}\}$

INITIATOR

Spontaneously

Unvisited $:= N(x)$

initiator $:= \text{true}$

VISIT

IDLE

receiving(T) // T token message

entry $:= \text{sender}$

Unvisited $:= N(x) - \text{sender}$

initiator $:= \text{false}$

VISIT

States $S = \{\text{INITIATOR}, \text{IDLE}, \text{VISITED}, \text{DONE}\}$

$S_{\text{init}} = \{\text{INITIATOR}, \text{IDLE}\}$

$S_{\text{term}} = \{\text{DONE}\}$

VISIT

```
if Unvisited  $\neq \emptyset$ 
    next  $\leftarrow$  Unvisited
    send(T) to next
    become VISITED
else
    if not initiator
        send(RET) to entry
    become DONE
```

VISITED

receiving(T) // T token message


Unvisited := Unvisited - sender
send(BEGDE) to sender


receiving(RET)


VISIT


receiving(BEDGE)


VISIT

































Complexity

Message Complexity:

Type of messages: token, back, return



$$2m = O(m)$$

$\Omega(m)$ is also a lower bound

Time Complexity:
(ideal time)

$$2m = O(m)$$

Totally sequential

$\Omega(n)$ is also a lower bound

Note:

most messages are on **Back Edges**


---> most time is spent on **Back Edges**

Idea: avoid sending messages on back edges


How ?

Using notification and **ACK** messages


DF+ Improving Time




DF+ Improving Time




DF+ Improving Time



DF+ Improving Time



DF+ Improving Time



DF+ Complexity

Message

Messages: Token, Return, Visited, Ack (ok)

Each entity (except init): receives 1 Token, sends 1 Return:
 $2(n-1)$

Each entity:

1 Visited to all neighbours except the sender

Let s be
the initiator

$$\sum |N(s)| + \sum_{x \neq s} (|N(x)| - 1)$$

$$= 2m - (n-1)$$

(same for Ack)

TOT: $4m$

DF+ Complexity

Time (ideal time)

Token and Return are sent sequentially: $2(n-1)$

Visited and Ack are done in parallel: $2n$

TOT: $4n - 2$




Summarizing:

DF Traversal

	Messages	Ideal Time
DF:	$2m$	$2m$
DF+:	$4m$	$4n - 2$


DF++

Do not send the Ack
What happens ?




DF++

Do not send the Ack
What happens ?




DF++

Do not send the Ack
What happens ?




DF++

Do not send the Ack
What happens ?




DF++

Do not send the Ack
What happens ?




DF++

Do not send the Ack
What happens ?



DF++

Do not send the Ack
What happens ?




A token is sent to an already visited node (= back edge)

Both nodes will eventually understand the "mistake"

DF++

Do not send the Ack
What happens ?



A token is sent to an already visited node (= back edge)

Both nodes will eventually understand the "mistake"

and pretend nothing happened

DF++ Complexity

In the worst case there is a "mistake" on each link except for the tree links

$$T \text{ \& Return} = 2n - 2$$

$$\text{Visited} = 2m - n + 1$$

$$\text{"Mistakes"} = 2(m - n + 1)$$

$$\text{Messages} = 4m - (n - 1)$$

BUT when we measure ideal time:

"mistakes" will not happen

$$\text{Time} = 2(n - 1)$$

Summary

	Messages	Ideal Time
DF:	$2m$	$2m$
DF+:	$4m$	$4n - 2$
DF++	$4m - n + 1$	$2n - 1$

Observations

An application:
access permission problems, e.g., Mutual Exclusion

Any Traversal does a Broadcast (not very efficient)
The reverse is not true.

Another Traversal: Smart Traversal

1- Build a Spanning Tree with SHOUT+

$$\text{Messages} = 2m$$

2- Perform DF Traversal

$$\text{Messages} = 2(n-1)$$

$$\text{Total Messages} = 2(m+n-1)$$

Another Traversal: Smart Traversal

1- Build a Spanning Tree with SHOUT+

Time $\leq d+1$ d : diameter

2- Perform DF Traversal


Time = $2(n-1)$

Total Time $\leq 2n+d-1$

Summary

	Messages	Ideal Time
DF:	$2m$	$2m$
DF+:	$4m$	$4n - 2$
DF++	$4m - n + 1$	$2n - 1$
Smart	$2m + 2n - 2$	$2n + d - 1$

Computations with Multiple initiator: WAKE-UP



FLOOD solves the problem.

General FLOOD algorithm: $O(m)$

More precisely: $2m - n + k^*$

↙
n. of initiators

1 init = broadcast = $2m - n + 1$

All init = $2m$

States $S = \{ASLEEP, AWAKE\}$

Sinit = {ASLEEP}

Sterm = {AWAKE}

Restrictions = R

ASLEEP

Spontaneously

send(W) to N(x)
become AWAKE

receiving(W)

send(W) to N(x) - sender
become AWAKE

Computations with Multiple initiator: WAKE-UP

In special topologies ?

TREE

Flood is optimal

$$n + k^* - 2$$

Computations with Multiple initiator: WAKE-UP

COMPLETE GRAPH

Broadcast

Flood
 $O(n^2)$

Specific
 $O(n)$

Wakeup

Flood

Specific
 $\Omega(n^2)$

Need additional assumptions
to reduce the complexity

HYPERCUBE

Broadcast

Flood
 $O(n \log n)$

Specific
 $O(n)$

Wakeup

Flood

Specific
 $\Omega(n \log n)$