

# Leader Election

---

## Chapter 3

Observations

Election in the Ring

Election in the Mesh


Election in an arbitrary graph

# Election



---

## The problem

Move the system from an initial configuration where all entities are in the same state into a final configuration where all entities are in the same state except one, namely the leader.




All entities are  
followers



The leader

# Election: Impossibility Result

---



Under the standard restrictions  $\mathbf{R}$  (BL, CN, TR)

## Theorem

The election problem cannot be generally solved if the entities do not have different identities.

# Election: Impossibility Result


---

Under the standard restrictions  $R$  (BL, CN, TR)

## Proof (Sketch)

By contradiction.

- Consider a system with two entities,  $x$  and  $y$ , in the same initial state
- If  $P$  solves the problem, it works in any conditions and delays
- Take a scenario synchronous and where  $x$  and  $y$  start simultaneously
- They have the same rules, thus they perform the same steps
- Thus, at each moment, they are doing the same thing
- If one becomes leader, even the other one (there cannot be a unique leader)



# Election: Additional Restriction

---

We need to extend the standard restrictions  $R$  (BL, CN, TR)

## The origin of the problem

Entities have all the same behaviour (behavioural symmetry)

- same set of rules
- same initial state

## The idea

Break the symmetry making each entity unique (in same way)

## Initial distinct values

Each entity has a unique identifier

# Election: Solution Strategy

---

## **Strategy elect minimum**

1. find the smallest value
2. the leader is the entity with that value

## **Strategy elect minimum initiator**



1. find the smallest value among the initiators
2. the leader is the entity with that value

## **Strategy elect root**

1. build a rooted spanning tree
2. the leader is the root of the tree

We can also consider taking the maximum

Note: with distinct Ids **Minimum Finding** is an election




# Election in the Tree

---

To each node  $x$  is associated a distinct identifier  $v(x)$

A simple algorithm:

- 1) Execute the saturation technique,
- 2) Choose the saturated node holding the minimum value





## **SATURATED**

*Receiving(Election, id\*)*

rules to add  
to Saturation

```
send("Termination") to N(x) - parent
if id(x) <= id* then
    become LEADER
else
    become FOLLOWER
```

## **PROCESSING**

*Receiving("Termination")*

```
send("Termination") to N(x) - parent
become FOLLOWER
```

## **Resolve**

```
send("Election", id(x)) to parent
become SATURATED
```

# Elect\_Min vs Elect\_Root in Trees

---

Complexity (n. messages)

Elect\_min:  $3n + k^* - 4$

Elect\_root:  $3n + k^* - 2$

Can we have a better estimate?

Bit Complexity


Elect\_min:  $n(c + \log id) + c(2n + k^* - 4)$

Elect\_root:  $2(c + \log id) + c(3n + k^* - 2)$

$c$  = bit of header

$\log id$  = bit of payload

# Ring



- $n$  entities
- $m = n$  links

- n. of entities = n. of links
- Symmetric topology
- Each entity has two neighbors



When there is sense of direction:



# Election Algorithms in Rings

---


- All the way
- As Far
- Controlled distance
- Electoral stages
  - bidirectional version
- Alternating steps

Electing the minimum

# All the way

---

**Basic Idea:** Each id is fully circulated in the ring.  
---> each entity sees all identities.



## ASSUMPTIONS


- Two versions: unidirectional/bidirectional links.
- Local orientation (i.e. not necessarily a sense of direction)
- Distinct identities.

# Basic Behaviour

---

1. Initiators start by selecting one neighbor and sending their ids

2. When a message arrives, an entity replies sending its id, forwarding the received message and storing the minimum id seen so far



## Correctness and Termination

---

The communication activities can terminate because entities do not forward messages with their ids.

**Question:** how can we make aware an entity of termination?

To terminate we need further assumptions:

- **FIFO assumption** (termination when an entity receives its id back)
- **until the number of messages reach the size of network** (knowledge of the size of the network)

Note: knowledge of  $n$  can be acquired

States:  $S = \{ASLEEP, AWAKE, FOLLOWER, LEADER\}$

$S_{INIT} = \{ASLEEP\};$

$S_{TERM} = \{FOLLOWER, LEADER\}.$

**IR;Ring**

**ASLEEP**

*Spontaneously*

**INITIALIZE**

**become AWAKE**

*Receiving(' `Election'', value, counter)*

**INITIALIZE;**

**send(' `Election'', value, counter+1)  
to other**

INITIALIZE

count:= 0

size:= 1

known:= false

**send("Election",id(x),size) to right;**

min:= id(x)

min:= Min{min, value}

count:= count+1

**become AWAKE**



## AWAKE

*Receiving ("Election", value, counter)*

If value  $\neq$  id(x) then

**send** ("Election", value, counter+1) **to other**

min:= MIN{min, value}

count:= count+1

if known = true then

**CHECK**

endif

else

ringsize:= counter

known:= true

**CHECK**

endif

**CHECK**

if count = ringsize then

if min = id(x) then

**become LEADER**

else

**become FOLLOWER**

endif

endif

## Complexity

---

Each identity crosses each link  $\rightarrow n^2$

The size of each message is  $\log(id + \text{counter})$

$O(n^2)$  messages

$O(n^2 \log(\text{MaxMsg}))$  bits


### Observations:

1. The algorithm also solves the data collection problem
2. It also works for unidirectional/bidirectional.

# AsFar (as it can)

---

**Basic Idea:** It is not necessary to send and receive messages with larger *ids* than the *ids* that have already been seen.





## ASSUMPTIONS

- Unidirectional/bidirectional ring
- Different *ids*
- Local orientation


# The Basic Idea

---



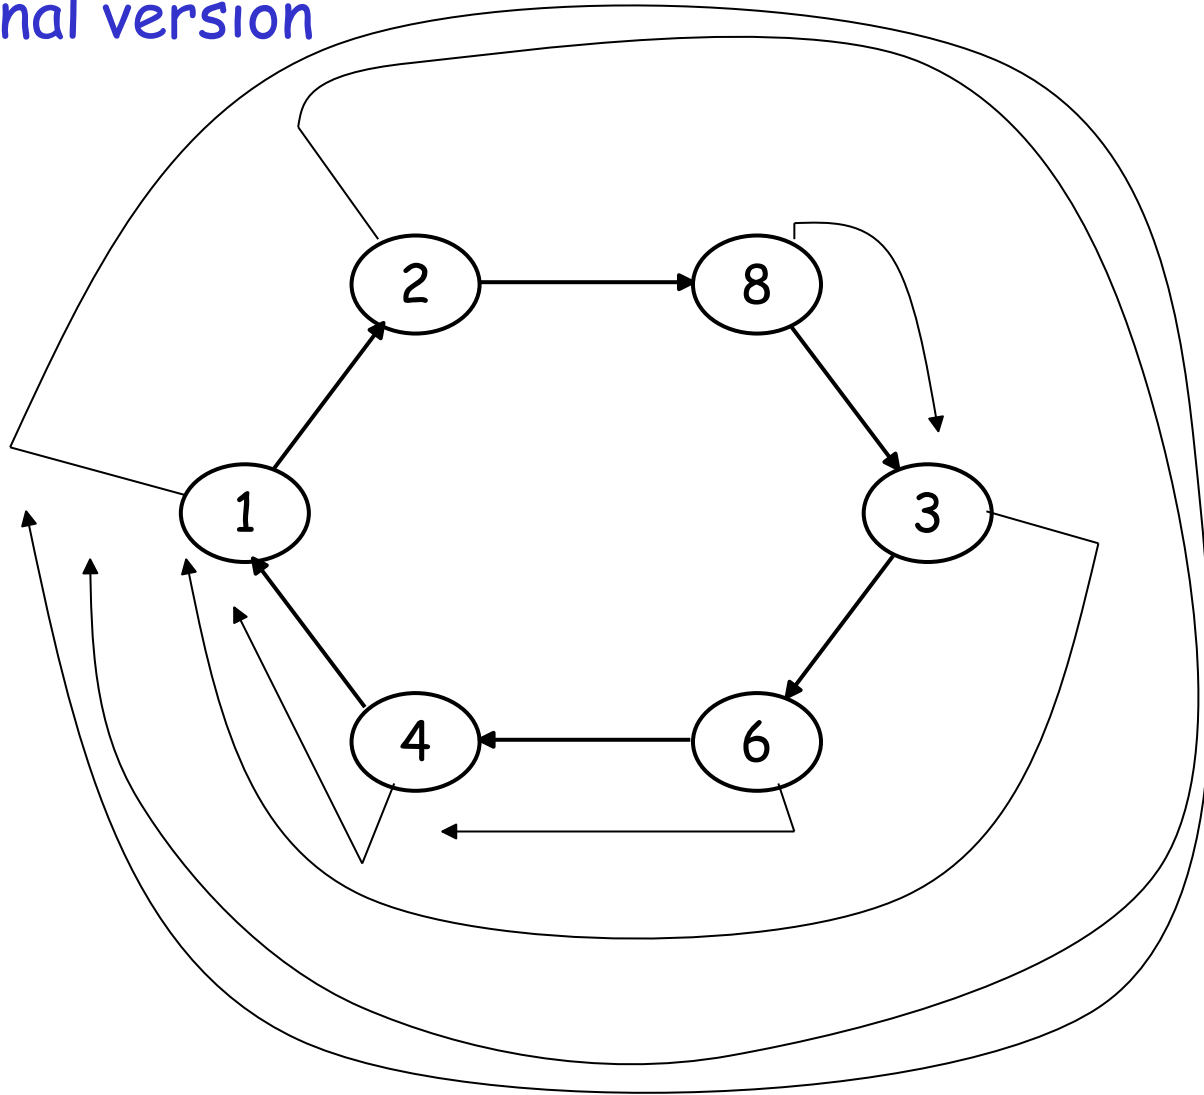


**Receiving  $y$  smaller-than me**  
**send( $y$ ) to other neighbour**




**Receiving  $y$  bigger-than me**  
**send( $x$ ) to other neighbour**  
(if not sent already)

unidirectional version



bidirectional version



States:  $S = \{ASLEEP, AWAKE, FOLLOWER, LEADER\}$

$S\_INIT = \{ASLEEP\}$

$S\_TERM = \{FOLLOWER, LEADER\}$

--- unidirectional version

## ASLEEP

*Spontaneously*

**send**("Election",id(x)) **to right**  
min:= id(x)  
**become** AWAKE

*Receiving("Election", value)*

**send**("Election",id(x)) **to right**  
min:= id(x)


If value < min then

**send**("Election", value) **to other**  
min:= value

endif

**become** AWAKE

**/\* this could be  
avoided if  
id(x) > value**





## AWAKE

*Receiving("Election", value)*

if value < min then

**send**("Election", value) **to other**

min:= value

else

If value = id(x) then NOTIFY endif

endif

*Receiving(Notify)*

**send**(Notify) **to other**

**become** FOLLOWER

NOTIFY

**send**(Notify) **to right**

**become** LEADER

## Correctness and Termination

---

The leader knows it is the leader when it receives its message back.

When do the others know ?


Notification is necessary !

### Observations:

- Bidirectional version

## Worst-Case Complexity (Unidirectional Version)

---



1 ---> n      links  
2 ---> n - 1    links  
3 ---> n - 2    links  
...      ...  
n ---> 1      link


$$n + (n - 1) + (n - 2) + \dots + 1 = \sum_{i=1}^n i = (n+1)(n) / 2$$

**Total:**  $n(n+1) / 2 + n = O(n^2)$

Last n: notification

## Best-Case Complexity (Unidirectional Version)

---



1 ---->  $n$  links  
for all  $i \neq 1$  ----> 1 link ( --> total =  $n - 1$  )

**Total:**  $n + (n - 1) + n = O(n)$

Last  $n$ : notification

## Average-Case Complexity

---

Entities are ordered in an equiprobable manner.

$$O(n \log n)$$

# Controlled Distance

---

Basic idea: Operate in stages. An entity maintains control on its own message.

## ASSUMPTIONS

- Bidirectional ring
- Different *ids*
- Local orientation

sense of direction only for simplicity - not needed

## Ingredients

---


1) Limited distance (to avoid big msgs to travel too much)

Ex: stage  $i$ : distance  $2^{i-1}$

2) Return messages (if seen something smaller does not continue)

3) Check both sides

4) Smallest always win (regardless of stage number)






---

*Candidate* entities begin the algorithm.

**Stage i:**

- Each *candidate* entity sends a message with its own *id* in both directions
- the msg will travel until it encounters a smaller Id or reaches a certain distance
- If a msg does not encounter a smaller Id, it will return back to the originator




- A candidate receiving its own msg back from both directions survives and start the next stage

Entities encountered along the path read the message and:

- Each entity  $i$  with a greater identity  $Id_i$  becomes defeated (passive).
- A defeated entity forwards the messages originating from other entities, if the message is a notification of termination, it terminates


**More...**

**Stage 1**




Distance = 1

**Stage 2**



Distance = 2






**Stage i**



Distance =  $2^{i-1}$

$2^{i-1}$

More...



States:  $S = \{\text{ASLEEP}, \text{CANDIDATE}, \text{DEFEATED}, \text{FOLLOWER}, \text{LEADER}\}$   
 $S\_INIT = \{\text{ASLEEP}\}$   
 $S\_TERM = \{\text{FOLLOWER}, \text{LEADER}\}$

## **ASLEEP**

*Spontaneously*

**INITIALIZE**

**become CANDIDATE**

*Receiving("Forth", id\*, stage\*, limit\*)*

**if** id\* < id(x) **then**

**PROCESS\_MESSAGE**

**become DEFEATED**

**else**

**INITIALIZE**

**become CANDIDATE**

## CANDIDATE

```
Receiving("Forth", id*, stage*, limit*)  
  if id* < id(x) then  
    PROCESS_MESSAGE  
    become DEFEATED  
  else  
    if id* = id(x) then NOTIFY
```

```
Receiving("Back", id*)  
  if id* = id(x) then CHECK
```

```
Receiving("Notify")  
  send("Notify") to other  
  become FOLLOWER
```

## DEFEATED

```
Receiving(*)  
  send(*) to other  
  if * = Notify then  
    become FOLLOWER
```

## INITIALIZE

```
stage := 1
limit := dis(stage)
count := 0
send("Forth", id(x), stage, limit) to N(x)
```

## PROCESS-MESSAGE

```
limit* := limit* - 1
if limit* = 0 then
    send("Back", id*, stage*) to sender
else
    send("Forth", id*, stage*, limit*) to other
```

## CHECK

```
count := count + 1
if count = 1 then
    count := 0
    stage := stage + 1
    limit := dis(stage)
    send("Forth", id(x), stage, limit) to N(x)
```

## NOTIFY

```
send("Notify") to right
become LEADER
```

## Correctness and Termination

---

If a candidate receives its message from the opposite side it sent it, it becomes the leader and notifies.

- The smallest id will always travel the max distance defeating every entity it encounters
- The distance monotonically increases eventually becoming greater than  $n$
- The leader will eventually receive its message from the opposite directions

**Note:** we do not need message ordering.

What happens if an entity receives a message from a higher stage ?



# Complexity

---

## Messages

When the distance is doubled at each stage i.e.,  $\text{dis}(i) = 2^{i-1}$ , the worst case complexity is

$$O(n \log n)$$

## Time (ideal)

- $2\text{dis}(i)$  time units required by the message with the smaller id to cover the distance
- $2n$  times unit are need to wake-ups & final notifications

$$2n + \sum_{i=1} 2 \text{dis}(i) = O(n)$$

# Stages

---

## Basic idea:


A message will travel until it reaches another candidate

A candidate will receive a message from both sides




## ASSUMPTIONS

- Distinct *ids*
- Bidirectional ring ( + unidirectional version)
- Local orientation
- **Message ordering** (for simplicity only: not needed)


Each *candidate* sends its own  $Id$  in both directions.



When a *candidate*  $i$  receives two messages  $Id_j$  (from the right) and  $Id_k$  (from the left), it determines if it becomes *passive* (= it is not the smallest), or if it remains *candidate* (= it is the smallest).



When a *candidate*  $i$  receives two messages  $Id_j$  (from the right) and  $Id_k$  (from the left),



After receiving the first message:  
**close-port** (enqueue messages possibly arriving later)

After receiving the second message, perform the action  
and **re-open-port**

States:  $S = \{\text{ASLEEP}, \text{CANDIDATE}, \text{WAITING}, \text{DEFEATED}, \text{FOLLOWER}, \text{LEADER}\}$

$S_{\text{INIT}} = \{\text{ASLEEP}\}$

$S_{\text{TERM}} = \{\text{FOLLOWER}, \text{LEADER}\}$

## **ASLEEP**

*Spontaneously*

**INITIALIZE**

**become CANDIDATE**

*Receiving("Election", id\*, stage\*)*

**INITIALIZE**

**min := MIN(id\*, min)**

**close-port(sender)**

**become WAITING**

## CANDIDATE

```
Receiving("Election", id*, stage*)  
  if id* <> id(x) then  
    min := MIN(id*, min)  
    close-port(sender)  
    become WAITING  
  else  
    send(Notify) to N(x)  
    become LEADER
```

## WAITING

```
Receiving("Election", id*, stage*)  
  open(other)  
  stage := stage + 1  
  min := MIN(id*, min)  
  if min = id(x) then  
    send("Election", id(x), stage) to N(x)  
    become CANDIDATE  
  else  
    become DEFEATED
```

## DEFEATED

*Receiving(\*)*

**send(\*)** to other

**if** \* = Notify **then**

**become** FOLLOWER

## INITIALIZE

stage := 1

count := 0

min := id(x)


**send**("Election", id(x), stage) to N(x)

## Correctness and termination

---

The minimal entity will never cease to send messages.

When an entity knows that it is the *leader*




it sends a *notification* message which travels around the ring.



## Complexity - Worst Case


At each step: At least half the entities became passive.

$$n_{i+1} \leq \frac{n_i}{2}$$



$$\begin{aligned} n_0 &= n \\ n_1 &= n/2 \end{aligned}$$

$$n_i = n/2^i$$



$$n/2^k = 1$$

when

$$k = \log n$$

**# steps:** At most  $(\log n)$

Each entity sends or resends 2 messages.


**# messages:**  $2n$

**# bits:**  $2n * (\log n)$

**Last entity:**  $2n$  messages to understand that it is the last active entity, then  $n$  notification messages.

**Total:**  $2n * (\log n) + 3n = O(n \log n)$

Best Case ?



## Removing Message Ordering

---

The ordering assumption ensures that messages received by a candidate in stage  $i$  are originated by candidates at the same stage


We can remove this assumption and modify the protocol to enforce that messages are processed in order

1. Each message carries the stage number of the entity originating it
2. When a candidate  $x$  in stage  $i$  receives a message  $M$  with stage  $j > i$ ,  $x$  stores  $M$  and processes it after all messages with stage  $i + 1, \dots, j - 1$  arrive

# Stages with Feedback

---

A feedback is sent back to the originator of the message




send YES to the smallest of the two IF it is smaller than me


A candidate  $x$  receives two messages from two neighboring candidates  $y=r(i,x)$  and  $z=l(i,x)$


A positive feedback is sent to

- $y$  if  $\text{id}(y) < \text{MIN}(\text{id}(x), \text{id}(z))$
- $z$  if  $\text{id}(z) < \text{MIN}(\text{id}(x), \text{id}(z))$



A candidate survives a stage if it receives positive feedback from its neighboring candidates





Candidates 8 and 5 do not receive any feedback

Candidate 2 is the only one that survives this stage

### Question

How can they do they didn't survive this stage?

### Answer

They do not step to the next stage, and become defeated when they receive the election message from 2 at the next stage (it works like a negative feedback).

## Correctness & Termination

---

Consider the entity with  $x_{\min}$  that will be the leader


It never sends a positive feedback and always receives two positive ones

This means that each neighboring candidate at each stage does not survive

The number of candidates at each stage is monotonically decreasing, until  $x_{\min}$  will be the only one

## Number of stages

If  $x$  survives at , it must have received a feedback from both neighbouring candidates  $l(i,x)$  and  $r(i,x)$   
Moreover,  $l^2(i,x)$  and  $r^2(i,x)$  do not survive too.



$$n_{i+1} \leq \frac{n_i}{3}$$

survivors at  
each stage



n. stages  
 $\log_3 n$




# Number of messages

---

At stage  $i$ , there are

- $2n$  election messages
- $n$  feedback messages

**Tot** =  $3n$  messages per stage



Messages for all stages  
 $3n + \log_3 n = O(n \log n)$

# Alternating Steps

---

Basic idea: Alternating directions.


- Different *ids*.
- Bidirectional ring and *sense of direction*.
- Local orientation.
- Message ordering.

send-left  
begin-to-defeat (if possible)  
send-right


# Algorithm

1. **Each entity** sends a message to its right. This message contains the entity's own *id*.
2. Each entity compares the *id* it received from its left to its own *id*.
3. If its own *id* is greater than the received *id*, the entity becomes passive.
4. All entities that remained active (surviving) send their *ids* to **their left**.
5. A surviving entity compares the *id* it received from its right with its own *id*.
6. If its own *id* is greater than the *id* it received, it becomes passive.
7. Go back to step 1 and repeat until an entity receives its own *id* and becomes *leader*.

Step 1: right



Step 2: left



States:  $S = \{\text{ASLEEP}, \text{CANDIDATE}, \text{WAITING}, \text{DEFEATED}, \text{FOLLOWER}, \text{LEADER}\}$

$S_{\text{INIT}} = \{\text{ASLEEP}\}$

$S_{\text{TERM}} = \{\text{FOLLOWER}, \text{LEADER}\}$

Restrictions: **IR**; Oriented Ring; Message Ordering

## **ASLEEP**

*Spontaneously*

**INITIALIZE**

**become CANDIDATE**

*Receiving("Election", id\*, step\*)*

**INITIALIZE**

**PROCESS\_MESSAGE**

**become CANDIDATE**

## CANDIDATE

*Receiving("Election", id\*, step\*)*  
    **if** id\*  $\neq$  id(x) **then**  
        PROCESS\_MESSAGE  
    **else**  
        **send**(Notify) to N(x)  
        **become** LEADER

## DEFEATED

*Receiving(\*)*  
    **send**(\*) to other  
    **if** \* = Notify **then**  
        **become** FOLLOWER

## INITIALIZE

step := step + 1  
min := id(x)  
**send**("Election", id(x), step) to right  
**close-port**(right)


## PROCESS\_MESSAGE

**if** id\* < min **then**  
        **open**(other)  
        **become** DEFEATED  
    **else**  
        step := step + 1  
        **send**("Election", id(x), step) to  
sender  
        **close-port**(sender)  
        **open**(other)

# Complexity


Analyze # of steps in worst case:

Last phase k



1 active entity


Phase k - 1



at least 2 active entities

(2) will become passive at the next step.


Phase k - 2



at least 3 active entities

(3) must be there; otherwise, (2) would be killed.

Phase k - 2



at least 5 active entities

1 2 3 5 8 13 21 ....



**# steps** = index of the lowest Fibonacci number  $\geq n$

$$F_1 = 1$$

$$F_2 = 2$$

$$F_3 = 3$$

$$F_4 = 5$$

$$F_5 = 8$$

...

$$F_k = i = ?$$

$$= \text{approx. } 1.45 \log_2 n$$

**# Messages** =  $n$  for each step

$$\text{Total} = \text{approx. } 1.45 n \log_2 n$$

## Conjecture:

In unidirectional rings,  
the worst case complexity is  $(n^2)$ ;  
to have a complexity of  $O(n \log n)$  messages,  
bidirectionality is necessary.

The Conjecture is false.

## Unidirectional version

---

**Simulation** of the bidirectional algorithm with the same complexity.


### **Example:** Unidirectional stages

Decompose the the operation of send in both directions in two steps

- Send the  $id(x)$  in the only possible direction
- Send also what you receive


The entities have exactly the same information as in the case of bidirectional links but shifted to the next candidate

Recall with bidirectional links



After sending the id

## Unidirectional Links



The same information but shift to next

# Election in Bidirectional Links (a recap)

---

<i>bidirectional</i>	<i>worst case</i>	<i>average</i>	<i>notes</i>
All the Way	$n^2$	$n^2$	
AsFar	$n^2$	$0.69n \log n + O(n)$	
ProbAsFar	$n^2$	$0.49n \log n + O(n)$	
Control	$6.31n \log n + O(n)$		
Stages	$2n \log n + O(n)$		
StagesFbk	$1.89n \log n + O(n)$		
Alternate	$1.44n \log n + O(n)$		oriented ring
BiMinMax	$1.44n \log n + O(n)$		
<i>lower bound</i>		$0.5n \log n + O(n)$	$n = 2^p$ known


# Election in Unidirectional Links

---

<i>unidirectional</i>	<i>worst case</i>	<i>average</i>	<i>notes</i>
All the Way	$n^2$	$n^2$	
AsFar	$n^2$	$0.69n \log n + O(n)$	
UniStages	$2n \log n + O(n)$		
UniAlternate	$1.44n \log n + O(n)$		
MinMax	$1.44n \log n + O(n)$		
MinMax+	$1.271n \log n + O(n)$		
<i>lower bound</i>		$0.69n \log n + O(n)$	
<i>lower bound</i>		$0.25n \log n + O(n)$	$n = 2^p$ known

# Mesh

---



If it is square mesh:  $n \text{ nodes} = n^{\frac{1}{2}} \times n^{\frac{1}{2}}$

$$m = O(n)$$

Asymmetric topology

corners

border

internal



**Idea:** Elect as a leader one of the four corners

Three phases:

1) Wake up

2) Election (on the border)  
among the corners



3) Notification

- Wake up

- Each initiator send a wake-up to its neighbours
- A non-initiator receiving a wake up, sends it to its other neighbours

$$O(m) = O(n)$$

## 2) Election on the border started by the corners



$O(n)$

### 3) Notification by flooding

$$O(m) = O(n)$$

$$\text{TOT: } O(n)$$